

Project 2

Route-Finding for an Unreliable Vehicle

CMSC 421, Fall 2019

Last update October 16, 2019

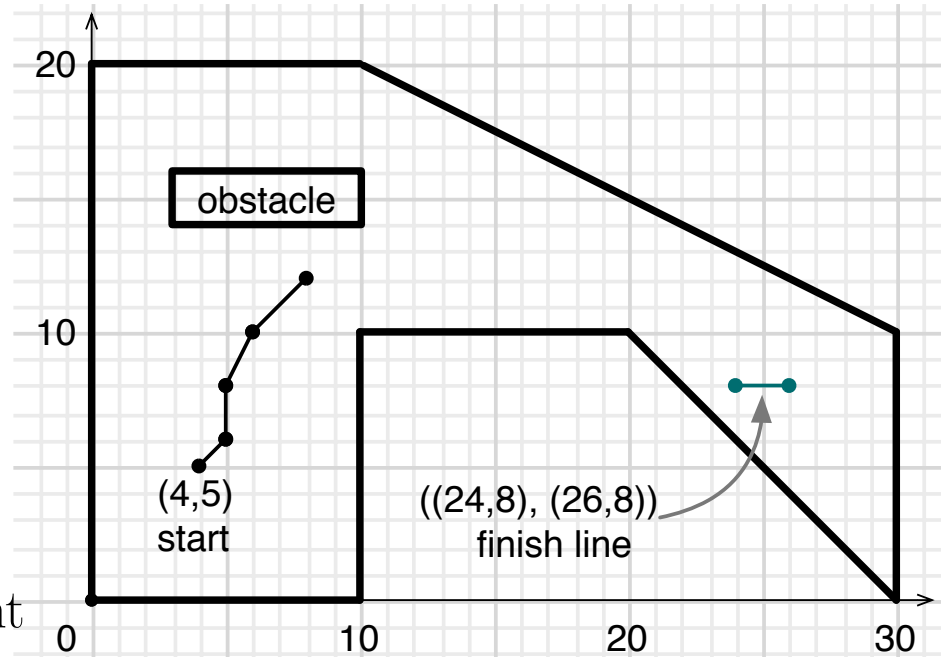
- ▶ Due date: November 1
- ▶ Late date (20% off): November 3

Another kind of racetrack problem

- Robot vehicle, starting point, finish line, walls are the same as in Project 1

Differences:

- (1) Vehicle's control system is unreliable
 - ▶ May move to a slightly different location than you intended
 - ▶ up to 1 unit in any direction
- (2) You can make bigger changes in velocity
 - ▶ Up to 2 units in any direction
- (3) Don't need to stop exactly on the finish line
 - ▶ OK to stop at distance ≤ 1



Moving the vehicle

- Current state $s = (p, z)$
 - location $p = (x, y)$, nonnegative integers
 - velocity $z = (u, v)$, integers

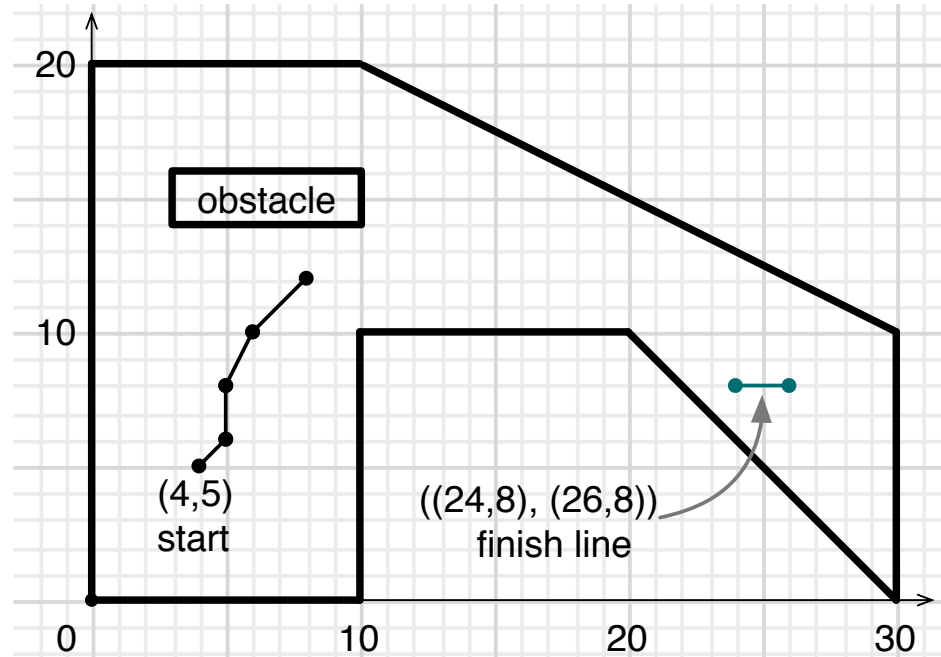
- You choose new velocity $z' = (u', v')$, where

$$u' \in \{u, u \pm 1, u \pm 2\},$$

$$v' \in \{v, v \pm 1, v \pm 2\}.$$

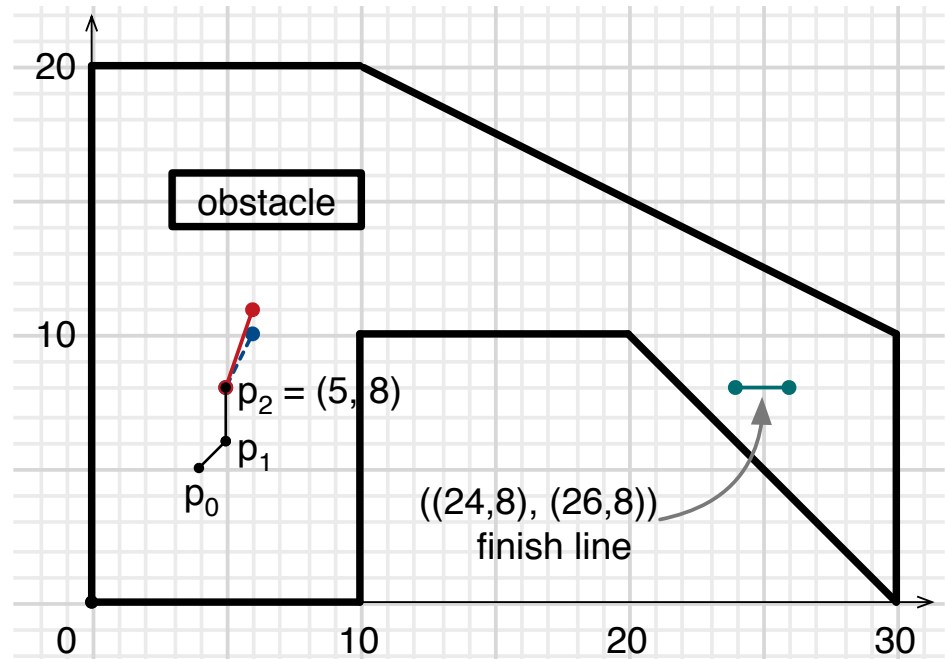
- If $z' \neq (0, 0)$, then the control system may make an error in your position
 - $e = (q, r)$, where $q, r \in \{-1, 0, 1\}$

- Vehicle moves to location $p' = p + z' + e = (x + u' + q, y + v' + r)$
- New state $s' = (p', z')$



Example

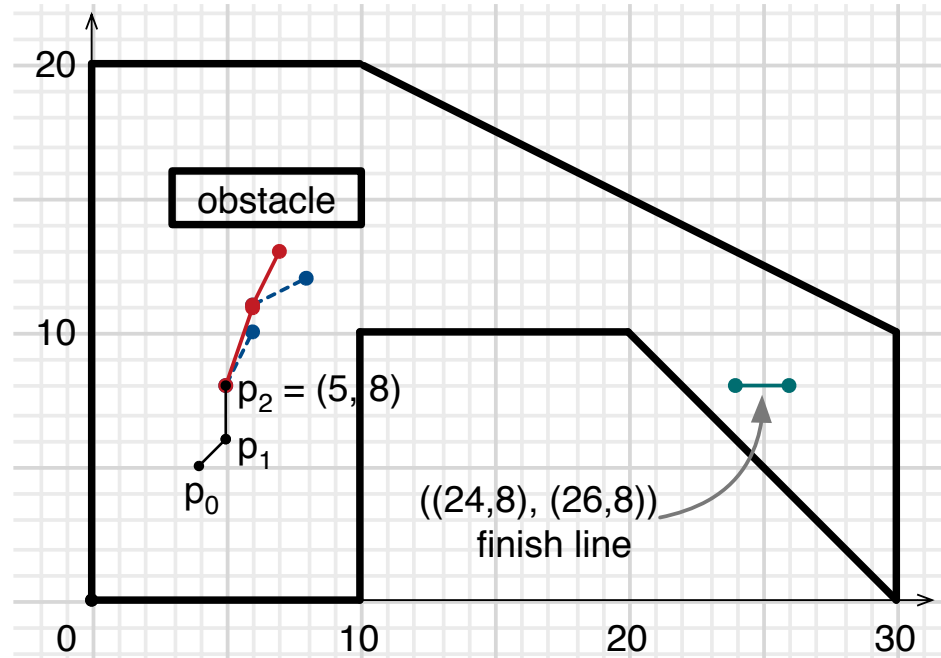
- State $s_2 = ((5, 8), (0, 2))$
 $p_2, \quad z_2$
- You choose
 $z_3 = z_2 + (1, 0) = (1, 2)$
- Control error $e_3 = (0, 1)$
- New location $p_3 = p_2 + z_3 + e_3$
 $= (5, 8) + (1, 2) + (0, 1)$
 $= (6, 11)$
- New state $s_3 = (p_3, z_3) = ((6, 11), (1, 2))$



- The control error doesn't change velocity, just your position
 - Unrealistic, but it makes the problems easier to solve

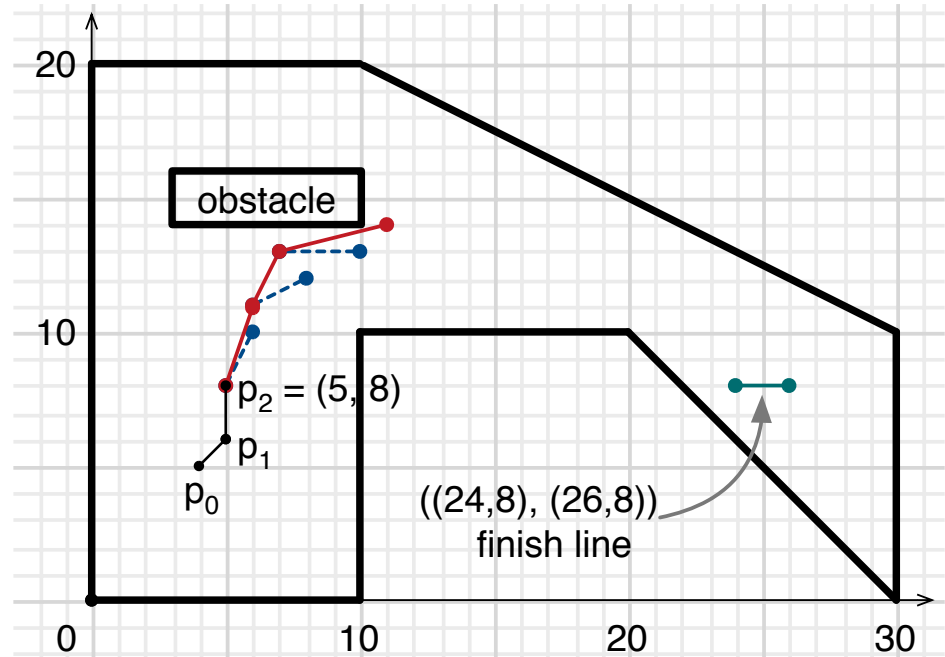
Example

- State $s_3 = ((6, 11), (1, 2))$
 $p_3, \quad z_3$
- You choose
 $z_4 = z_3 + (1, -1) = (2, 1)$
- Control error $e_4 = (-1, 1)$
- New location $p_4 = p_3 + z_4 + e_4$
 $= (6, 11) + (2, 1) + (-1, 1)$
 $= (7, 13)$
- New state $s_4 = (p_4, z_4) = ((7, 13), (2, 1))$



Example

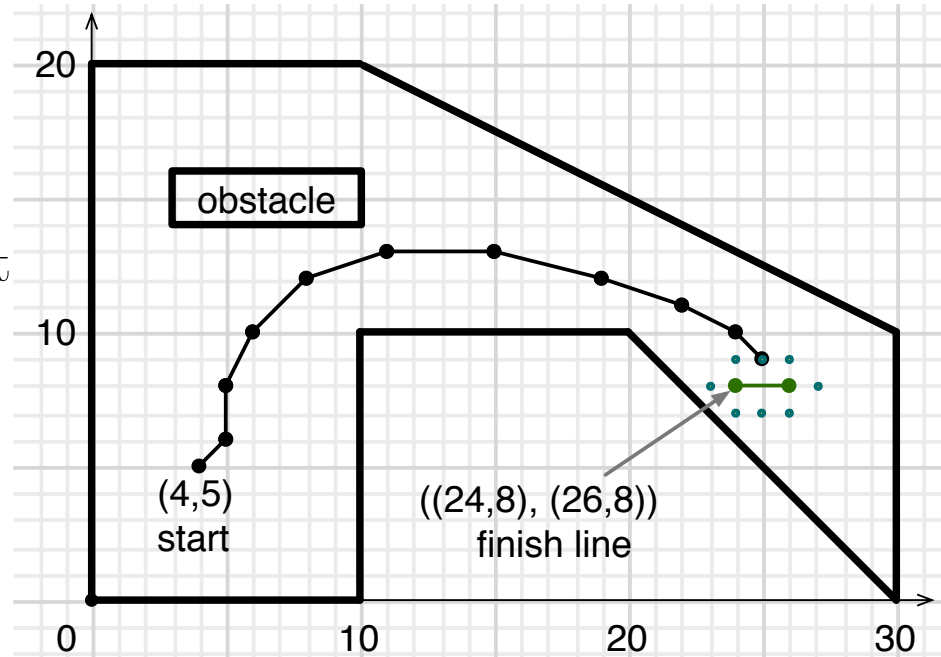
- State $s_4 = ((7, 13), (2, 1))$
 $p_4, \quad z_4$
- You choose
 $z_5 = z_4 + (1, -1) = (3, 0)$
- Control error $e_5 = (1, 1)$
- New location $p_5 = p_4 + z_5 + e_5$
 $= (7, 13) + (3, 0) + (1, 1)$
 $= (11, 14)$



- New state $s_5 = (p_5, z_5) = ((11, 14), (3, 0))$
- Trajectory is *unsafe*
 - Would have crashed if e_5 were $(0, 1)$ or $(-1, 1)$
- Ideally, you want a strategy that will always keep you from crashing regardless of what control errors occur

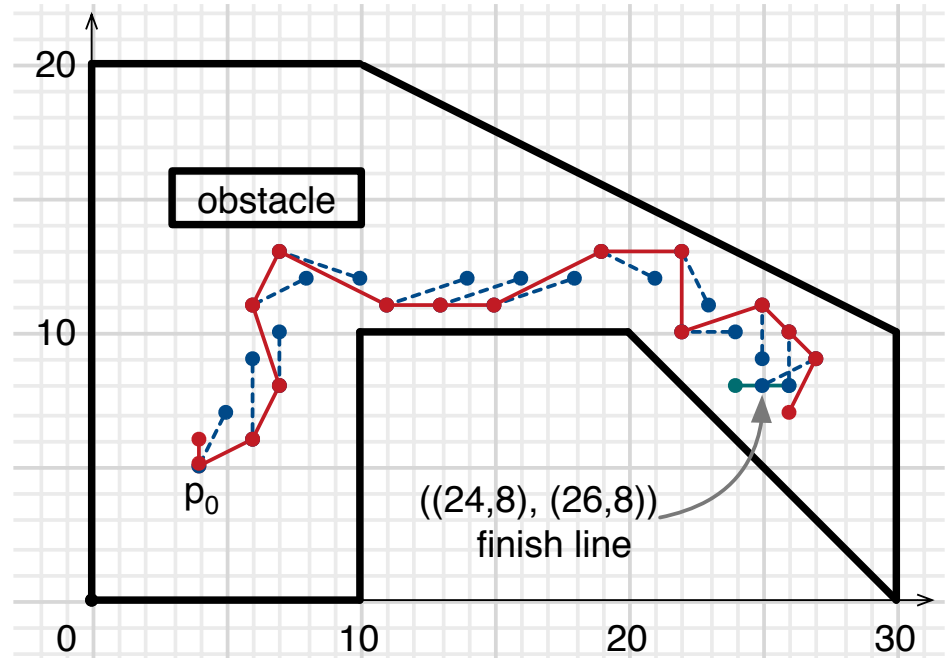
Objective

- Get to the finish line and stop
 - ▶ Might not be able to land exactly on the line
 - ▶ Control errors can prevent that
- OK to get to distance ≤ 1
- Need to stop
 - ▶ Last move needs to have velocity 0, as in Project 1
- Want to get there as quickly as possible without crashing, despite control errors



Strategy

- Pretend the control system is an opponent that's trying to make you crash
- Choose moves that will keep you from crashing, regardless of what it does
- Write a game-playing algorithm to do it move by move
 - as in chess, checkers, or go



How to do it

- One possibility: alpha-beta game-tree search
 - ▶ Limited-depth search, static evaluation function
- Another possibility: Monte Carlo rollouts
 - ▶ Problem: randomly generated paths are very unlikely to go to the goal
 - ▶ I don't think it will work very well
- Another idea: biased Monte Carlo rollouts
 - ▶ Generate paths randomly, but bias the moves toward good evaluation-function values
 - ▶ How well this will work, I have no idea
- No way to guarantee you won't crash

Opponent

- We'll give you a simple opponent program
 - ▶ It will try to make you crash, but won't be very intelligent about it
- Warning: don't write a program that just tries to take advantage of the dumb opponent!
 - ▶ When we grade your program, we'll use a more intelligent opponent
- Need to choose moves that won't crash, no matter what the opponent does

Other comments

- You may use any of the code I gave you for Project 1, and any of the code you developed for Project 1
 - ▶ You can modify it if you wish
- Caveat: most of it won't be very useful
 - ▶ You'll need to write a game-tree-search algorithm and/or a Monte Carlo rollout algorithm
- You'll need a heuristic function
 - ▶ You can use the one you developed for Project 1
 - ▶ You can use any of the ones I gave you for Project 1
 - e.g., `h_walldist`
- Caveat: Will a heuristic function for Project 1 work well as a game-tree-search heuristic?
 - ▶ You might need to make modifications

What you need to submit

- You need to submit a file called `proj2.py` containing a program called `main`
- We'll give you a game environment for running it
 - ▶ It will simulate turn-by-turn interactions with the opponent
 - ▶ At each turn, it will run `proj2.main(s, f, w)`
 - s = state, f = finish line, w = list of walls
- Your `proj2.main` program should print (to standard output) a sequence of choices for what velocity to use. Each choice should be a pair of integers (u, v) followed by a linebreak.
 - (2, 2)
 - (1, 3)
 - (1, 2)
 - (1, 2)
- Keep searching for better and better recommendations
 - e.g., iterative deepening, or additional Monte Carlo rollouts

More about the game environment

- Game environment runs your `proj2.main` program as a separate process
 - ▶ Lets it run for 5 seconds, kills it, reads the last velocity it chose
- After getting your chosen velocity (u, v) , it lets the opponent choose what error to use
 - ▶ $e = (q, r)$, where $q, r \in \{-1, 0, 1\}$
- It computes the new state, and checks whether the game has ended
 - ▶ you crash \Rightarrow you lose
 - ▶ you reach the finish line and your velocity is $(0,0) \Rightarrow$ you win
 - ▶ otherwise, game hasn't ended \Rightarrow game environment will call your program again, with the new current state
- If the game hasn't ended, it goes to the next turn
 - ▶ runs your `proj2.main` program again

Files I'll provide

- File on Piazza: `project2b_code.zip`
 - ▶ Not `project2_code.zip` – that version had an error in it
- `sample_probs` – modified version of the test problems from Project 1.
 - ▶ I removed or modified the ones that were obviously unsolvable.
 - ▶ Each problem is a list of the form `[name, p0, finish, walls]`
 - ▶ If a problem's dimensions are so small that the problem is unsolvable, you can call `double(p)` where p is the problem, to return a problem in which the x and y dimensions are both doubled.
- `opponents.py` – two simple opponent programs.
 - ▶ `opponent1` tries to head for the wall
 - ▶ `opponent0` makes moves at random
 - ▶ By default, `env.main` uses `opponent0`

Files I'll provide (continued)

- `env.py` – environment for running your `proj2.py` file.

Here's what `env.main(problem)` does:

1. If you have a `proj2.initialize`, it launches `proj2.initialize(s, f, w)`, waits 5 seconds, and kills the process if it hasn't exited.
 - ▶ This is so you can compute some data to use in your `proj2.main` program
 - ▶ Your `proj2.initialize` should write the data to a file called `data.txt`; otherwise the data will be lost when the process exits
2. It repeats the following steps until you win or lose:
 - ▶ Launch `proj2.main`, wait 5 seconds, and kill the process
 - ▶ Read the last velocity (u, v) in `choices.txt`. If it isn't legal, return `lose`
 - ▶ If $(u, v) = (0, 0)$ and distance from finish line ≤ 1 , return `win`.
 - ▶ Call the opponent to add an error to the velocity
 - ▶ Draw the move using turtle graphics
 - ▶ If the move crashes into a wall, return `lose`

Files I'll provide (continued)

- `proj2_example.py` – a deliberately stupid version of `proj2.py`
 - ▶ Rename it to `proj2.py` if you want to use it with `env.py`.
 - ▶ I provided it so you can see how `env.py` works before you start writing your own `proj2.py`.
 - ▶ It can win if it's lucky, but usually it eventually crashes
 - ▶ You'll need to write something that works much better
- Some files used by `proj2_example.py`
 - ▶ `racetrack_example.py` – modified version of `racetrack` from Project 1
 - ▶ `fsearch.py` and `tdraw.py` – same as in Project 1

Grading

- Evaluation criteria:
 - ▶ 35% correctness: – whether your algorithm works correctly, whether your submission follows the instructions
 - ▶ 15% programming style – see the following
 - Style guide: <https://www.python.org/dev/peps/pep-0008/>
 - Python essays: <https://www.python.org/doc/essays/>
 - ▶ 15% documentation
 - Docstrings at start of the file and in each function; comments elsewhere
 - ▶ 35% on performance
 - Does your program crash? If so, then how frequently?
 - If it doesn't crash, then how many moves to reach the finish line?
 - Top n performers ($n \approx 3$ to 5) will get extra credit